

# Supplementary Material: Auditing Annotation Support and Evaluation Sensitivity in Aerial Vehicle Detection

Xinling Wen, Jiahao Zhang, Wenqi Liu, Yangxu Ren, and Yu Chen

## S1. NOTATION AND AUDIT QUANTITIES

This supplement specifies the formal objects and the data-free records underlying the manuscript. All reported quantities are restricted to the declared artifacts, mapping, common-coverage denominator, frozen candidate pool, and registered rule set. They are not estimates of annotation-error rates, a general detector leaderboard, or physical traffic density.

Let  $G_s$  be the valid mapped ground-truth boxes for source  $s$ , and let  $G_s(r) \subseteq G_s$  be the boxes retained by a support rule  $r$ . The annotation-support shift is

$$S_s(r) = 1 - \frac{|G_s(r)|}{|G_s|}.$$

For candidate  $k$  and metric  $m$ , the fixed- $(d, \tau, c, \iota)$  policy band is  $B_m^\pi = \max_{r \in \Pi_s} m(d, \tau, \pi, c, \iota) - \min_{r \in \Pi_s} m(d, \tau, \pi, c, \iota)$ . The authoritative primary policy set is  $\Pi = \{I, O, N\}$ : include-all ( $I$ ), exclude/source-ignore-off ( $O$ ), and exclude/source-ignore-on ( $N$ ). The ancillary “exclude-and-ignore-removed” control is intentionally outside this set and therefore does not enter bands, headline record counts, or Fig. 2 of the main paper.

For a finite headline rule set  $\mathcal{R}_h$ , define the candidate-level rule band  $B_k(\mathcal{R}_h) = \max_{r \in \mathcal{R}_h} m_k(r) - \min_{r \in \mathcal{R}_h} m_k(r)$ . For a reference rule  $r_0 \in \mathcal{R}_h$ , winner  $k^*$ , and competitor  $j$ , define  $\Delta_j(r) = m_{k^*}(r) - m_j(r)$  and

$$\Gamma_j(\mathcal{R}_h) = \max_{r \in \mathcal{R}_h} |\Delta_j(r) - \Delta_j(r_0)|.$$

$\Gamma_j(\mathcal{R}_h)$  is the differential policy radius. It removes common-mode score movement and is the relevant deterministic diagnostic for a pairwise rank claim. We additionally report  $\min_{r \in \mathcal{R}_h} \Delta_j(r)$ , the direct observed rank margin, and a paired-bootstrap winner probability for the declared resampling unit.

The main paper’s rule tuple is  $r = (d, \tau, \pi, c, \iota)$ , where  $d$  selects absolute or normalized support,  $\tau$  is the support threshold,  $\pi$  is a small-object policy,  $c$  is confidence, and  $\iota$  is IoU. AP integrates over score order and hence has no fixed confidence coordinate. All multi-candidate comparisons are computed on  $C_s$ , defined before class filtering from raw-cache image names. A covered image with no retained vehicle prediction remains an evaluated empty prediction.

## S2. FORMAL RESULTS AND PROOFS

**Lemma S1 (resize invariance).** Under isotropic resizing by factor  $\alpha > 0$ ,  $(w, h, W, H)$  maps to  $(\alpha w, \alpha h, \alpha W, \alpha H)$ . Thus

TABLE I  
NOTATION USED IN THE SUPPLEMENTARY ANALYSIS.

| Symbol                    | Definition                                            |
|---------------------------|-------------------------------------------------------|
| $G_s, G_s(r)$             | Valid mapped boxes; boxes retained by rule $r$        |
| $S_s(r)$                  | Annotation-support shift $1 -  G_s(r) / G_s $         |
| $m_k(r), B_m^\pi$         | Candidate metric; fixed-condition policy band         |
| $B_k(\mathcal{R}_h)$      | Candidate rule band over the headline rule set        |
| $\Delta_j(r)$             | Winner-competitor margin under rule $r$               |
| $\Gamma_j(\mathcal{R}_h)$ | Maximum margin change over $\mathcal{R}_h$ from $r_0$ |
| $C_s$                     | Source-wise common image-coverage denominator         |
| $T, P, N$                 | Global TP, FP, and FN counts for micro- $F_1$         |
| $D$                       | $2T + P + N$ , the micro- $F_1$ denominator           |

$s'_a = \max(\alpha w, \alpha h) = \alpha s_a$ , while

$$s'_n = \max\left(\frac{\alpha w}{\alpha W}, \frac{\alpha h}{\alpha H}\right) = s_n.$$

Consequently, a fixed-pixel threshold depends on image scale, whereas a normalized-side threshold is invariant to isotropic image resizing. This is an invariance statement about the audit quantity, not a claim that normalized filtering removes every source difference.

**Proposition S1 (policy-robust candidate-pool winner).** Suppose  $k^*$  is the reference winner and  $\Delta_j(r_0) > 0$  for every competitor  $j$ . If  $\Delta_j(r_0) > \Gamma_j(\mathcal{R}_h)$  for all  $j$ , then  $k^*$  is the winner for every  $r \in \mathcal{R}_h$ .

*Proof.* By the definition of  $\Gamma_j(\mathcal{R}_h)$ ,

$$\Delta_j(r) \geq \Delta_j(r_0) - |\Delta_j(r) - \Delta_j(r_0)| \geq \Delta_j(r_0) - \Gamma_j(\mathcal{R}_h) > 0.$$

Thus  $m_{k^*}(r) > m_j(r)$  for every competitor and every registered rule.  $\square$

**Corollary S1.** Since

$$\Gamma_j(\mathcal{R}_h) \leq B_{k^*}(\mathcal{R}_h) + B_j(\mathcal{R}_h),$$

the sufficient condition  $\Delta_j(r_0) > B_{k^*}(\mathcal{R}_h) + B_j(\mathcal{R}_h)$  guarantees the same conclusion.

*Proof.* Write  $x_k(r) = m_k(r) - m_k(r_0)$ . Each  $x_k(r)$  lies between the minimum and maximum movement of candidate  $k$ , whose range over  $\mathcal{R}_h$  is  $B_k(\mathcal{R}_h)$ . Hence  $|x_{k^*}(r) - x_j(r)| \leq B_{k^*}(\mathcal{R}_h) + B_j(\mathcal{R}_h)$ , and maximization over  $r \in \mathcal{R}_h$  gives the inequality.  $\square$

The corollary is deliberately conservative: it ignores correlation in rule-induced movements. For example, let  $(m_A(r_0), m_B(r_0)) = (.70, .68)$  and, under another rule,  $(.60, .58)$ . Both individual bands equal .10, exceeding the

TABLE II  
PAIRWISE DIAGNOSTICS OVER SIX NOMINAL SCALE-POLICY SETTINGS.  
EVERY ROW HAS A DETERMINISTIC CERTIFICATE; BOOTSTRAP  
QUALIFICATION IS REPORTED SEPARATELY.

| Source      | Metric           | $\Delta_0$ | $\Gamma$ | min $\Delta$ | Cert. | Bootstrap     |
|-------------|------------------|------------|----------|--------------|-------|---------------|
| VisDrone    | AP <sub>50</sub> | .3544      | .0206    | .3544        | yes   | all > 0       |
| Local UAVDT | AP <sub>50</sub> | .0340      | .0326    | .0014        | yes   | all include 0 |
| VisDrone    | $F_1$            | .3354      | .0026    | .3354        | yes   | all > 0       |
| Local UAVDT | $F_1$            | .1911      | .00007   | .1910        | yes   | all include 0 |

reference margin .02, yet  $\Delta_B(r) = .02$  throughout and  $\Gamma_B = 0$ . Conversely, anti-aligned movements can make small individual bands consequential for a close race. Differential radii therefore complement, rather than replace, the individual policy bands.

### S3. COVERAGE IDENTIFIABILITY AND AN $F_1$ BOUND

**Proposition S2 (non-identifiability of full-population micro- $F_1$  ranking under unequal coverage).** Let candidates  $A$  and  $B$  be observed on different image subsets  $C_A$  and  $C_B$ . If their TP, FP, and FN contributions on unobserved images are unconstrained, the ordering of their full-population micro- $F_1$  scores is not identifiable from the observed subset counts.

*Proof.* Keep all observed counts fixed. On unobserved images, construct one completion with sufficiently favorable TP/FP/FN contributions for  $A$  and unfavorable contributions for  $B$ ; it produces an  $A$ -leading full-population micro- $F_1$ . Reverse these unconstrained contributions to produce a  $B$ -leading completion with exactly the same observed subset counts. Thus no unique full-population micro- $F_1$  ordering is identified.  $\square$

An analogous construction applies to AP when scores and detections on unobserved images are unrestricted; it is not needed for the formal micro- $F_1$  result above.

This result motivates the source-wise common denominator used for VisDrone and local UAVDT. It also gives the formal boundary for AI-TOD: its 1,869 covered images have 16,900 vehicle boxes (9.0 per image; median side 12 px), while 935 noncovered images have 43,004 boxes (46.0 per image; median side 16 px). The observed coverage-conditioned sensitivity analysis cannot be extrapolated into a full-source ranking without missing-coverage assumptions.

**Proposition S3 (count-perturbation bound for micro- $F_1$ ).** For  $F(T, P, N) = 2T/(2T + P + N)$ , let  $D_{\min} > 0$  lower-bound  $D = 2T + P + N$  along the segment between two nonnegative count vectors. Then

$$|\Delta F| \leq \frac{2|\Delta T| + |\Delta P| + |\Delta N|}{D_{\min}}.$$

*Proof.* The partial derivatives satisfy  $|\partial F/\partial T| = 2(P + N)/D^2 \leq 2/D$ , and  $|\partial F/\partial P| = |\partial F/\partial N| = 2T/D^2 \leq 1/D$ . Apply the integral form of the mean-value bound along the line segment and use  $D \geq D_{\min}$ .  $\square$

The bound does not predict an observed band from support removal alone, because the signs and magnitudes of TP, FP, and FN changes remain detector- and policy-dependent. It does explain why extreme removal permits a large band: at 24 px, the pooled excluded fractions are 26.7%, 4.9%, and 90.3% for VisDrone, local UAVDT, and AI-TOD, respectively. AI-TOD's .330  $F_1$  band is therefore interpreted jointly with

TABLE III  
ARTIFACT-LEVEL PROVENANCE AND EVALUATION SCOPE.

| Source      | Eval. imgs | Covered | Vehicle boxes | Mapping             | Scope            |
|-------------|------------|---------|---------------|---------------------|------------------|
| VisDrone    | 548        | 548     | 17,040        | car/van/truck/bus   | full audit       |
| Local UAVDT | 305        | 305     | 11,054        | car/truck/bus       | full audit       |
| AI-TOD      | 2,804      | 1,869   | 59,904        | vehicle             | sensitivity only |
| DOTA-v2     | 5,108      | –       | 53,349        | small/large vehicle | structural only  |

retained support and removed-object treatment, not as an isolated detector effect.

### S4. ARTIFACT PROVENANCE AND COVERAGE AUDIT

Table III records the source-specific boundaries that matter for interpretation. DOTA-v2 contributes only structural support: its quadrilaterals are converted to axis-aligned horizontal-box proxies for the scale audit, so its distances are representation-conditioned. The local UAVDT artifact is a COCO-style conversion and is conversion-conditioned. AI-TOD's available crop reconstruction is xView-derived and has incomplete prediction coverage; it is not used for a rank claim.

The vehicle-positive image counts are 519/548 for VisDrone, 305/305 for local UAVDT, 1,443/2,804 for AI-TOD, and 2,581/5,108 for DOTA-v2. Median horizontal-box maximum sides are 41, 33, 14, and 18 px, respectively; median normalized sides are .0412, .0278, .0175, and .0352. These are descriptive support properties. They do not establish which source is "more correct," nor do they assert a causal resolution effect.

Raw prediction coverage is determined by image-name presence before mapping predictions into the vehicle superclass. Consequently, a covered image whose predictions are entirely filtered out still contributes all valid mapped ground truth as false negatives, rather than vanishing from AP or  $F_1$ . Source ignores are preserved for VisDrone. Its truncation and occlusion attributes are not filtered. For local UAVDT, the applied coordinates are the recorded  $xywh$  conversion. The provenance manifest in the repository supplies raw paths and hashes without redistributing licensed images or prediction caches.

The published package at <https://anonymous.4open.science/t/aerial-evaluation-audit-5565/README.md> contains the ontology mapping, candidate manifest, coverage records, corrected headline outputs, matching controls, bootstrap records, figure-source summaries, scripts, and SHA-256 manifest. It is sufficient to inspect every claim-scoped derived record; acquiring the licensed source artifacts remains the responsibility of the reproducer.

### S5. FULL RULE CONTRACT AND MATCHING CONTROLS

The full greedy surface comprises seven candidates, eight confidence values, four IoUs, eight absolute or six normalized support thresholds, and three primary policies:  $7[8 \cdot 4 \cdot 8 \cdot 3 + 8 \cdot 4 \cdot 6 \cdot 3] = 9,408$  records. The absolute thresholds are {16, 20, 24, 28, 32, 40, 48, 64} px, normalized thresholds are {.005, .0075, .010, .015, .020, .030}, confidence values are {.10, .15, .20, .25, .30, .35, .40, .50}, and IoUs are {.25, .30, .40, .50}. Headline rules use 24 px and .015, registered interior grid values selected before headline computation.

TABLE IV  
POLICY SEMANTICS AND ROLE IN REPORTED QUANTITIES.

| Policy                     | Removed valid boxes                      | Use                  |
|----------------------------|------------------------------------------|----------------------|
| Include-all                | retained                                 | primary grid         |
| Exclude/source-ignore-off  | absent; overlaps may be FP               | primary grid         |
| Exclude/source-ignore-on   | absent; source ignores suppress overlaps | primary grid         |
| Exclude-and-ignore-removed | absent; removed overlaps suppressed      | ancillary diagnostic |

For fixed-operation  $F_1$ , predictions are score ordered. Each prediction matches the highest-IoU eligible yet-unmatched valid ground-truth box above the declared IoU; unmatched predictions are FP and unmatched retained boxes are FN. The Hungarian control first removes invalid edges, maximizes cardinality of valid matches, and then maximizes total IoU among maximum-cardinality solutions. Dummy nodes permit unmatched detections and boxes. Maximizing cardinality before total IoU prevents an IoU-optimal assignment from reducing the number of valid matches.

AP uses one-category COCO bbox evaluation with complete score order, IoU .5 for AP<sub>50</sub>, standard area ranges, and registered maxDets = 100. The 100-detection cap is material: after source-ignore handling, it truncates predictions in 406–490 VisDrone images and 272–292 local-UAVDT images across primary policies. Focused 300 and 2000 controls preserve every one of the 10 nonduplicate multi-candidate AP winners; at 2000 no retained prediction is truncated (the largest retained count is 1,066), but AP<sub>50</sub> pairwise gaps relative to 100 shift by [−.041, .048].

The ancillary control is especially informative for AI-TOD: under it,  $F_1$  is .130 at 24 px and .317 at .015, compared with .065 and .260 under exclusion without protection. Thus the observed policy variation combines extreme support removal and the treatment of predictions that overlap removed targets.

#### S6. ADDITIONAL RULE-SURFACE DIAGNOSTICS

The main paper intentionally reports two compact figures. This section supplies the full-headline-rule visualizations used for cross-checking the policy summaries. The surfaces below are derived only from the registered grids; they introduce neither new candidates nor new evaluation rules. Their purpose is diagnostic: they make clear that a stable winner is a property of the declared pool under common coverage, rather than a claim about a detector family.

The absolute and normalized surfaces should not be treated as a cross-source leaderboard. They are useful because they expose where the evaluation target changes: a pixel threshold moves under resizing, while a normalized-side threshold remains invariant under isotropic resizing. The retained-support analysis in the main paper shows why the same numerical rule cannot be interpreted as a source-neutral target definition.

#### S7. COMPLETE HEADLINE BOOTSTRAP RECORDS

Tables V–VI make the uncertainty qualification explicit. Each row is a separate paired bootstrap interval; endpoint ranges in the main paper are summaries across these rows, not a single pooled confidence interval. The local-UAVDT rows use the filename-derived sequence proxy described in the main paper and Section S8.

The apparent duplication in some source-ignore rows is itself an audit finding. Source-ignore-on and source-ignore-off are identical for local UAVDT under the registered artifacts, so four nominal AP outcomes collapse to fewer distinct policy semantics. We retain nominal row-level records for transparent replay and separately state the 10 nonduplicate multi-candidate AP outcomes and 20 nonduplicate matching outcomes in the main paper.

#### S8. BOOTSTRAP, REPRODUCIBILITY, AND CLAIM AUDIT

Pairwise uncertainty resamples the common evaluation unit jointly for the leading and runner-up candidates. VisDrone uses image units. Local UAVDT uses a filename-derived sequence proxy (substring before the first underscore): 93 units over 305 images, with min/median/max unit sizes of 1/1/104. Each analysis uses 1,000 paired replicates, seed 20260712, and percentile 2.5/97.5% intervals. Replicates recompute global micro- $F_1$  from resampled TP/FP/FN totals and recompute score-ordered one-category COCO AP<sub>50</sub> at maxDets = 100.

For AP<sub>50</sub>, all six VisDrone policy-specific intervals favor SAHI: their lower endpoints range from .339 to .362 and their upper endpoints from .366 to .387. Every one of the four nonduplicate local-UAVDT sequence-proxy intervals contains zero, with lower endpoints from −.310 to −.296, upper endpoints .071–.111, and winner probabilities .471–.604. For  $F_1$ , all 12 VisDrone intervals are positive with winner probability 1.00, whereas all 12 local-UAVDT sequence-proxy intervals contain zero despite deterministic tiling wins (winner probabilities .691–.811). The latter unit is a sensitivity device, not evidence that filename prefixes recover official UAVDT sequence identities.

#### Reproducibility and claim-audit contract.

| Item                       | Released or fixed record                     |
|----------------------------|----------------------------------------------|
| Rule contract              | configs/locked_rule_grid.yaml                |
| Ontology/mapping           | label_reliability/config/ontology.json       |
| Coverage denominator       | outputs/coverage/*.csv                       |
| AP/ $F_1$ headline records | outputs/ap_headline, coverage_corrected_grid |
| Matcher control            | outputs/matching_headline/*.csv              |
| Bootstrap outputs          | outputs/statistics/bootstrap_headline/*.csv  |
| Policy diagnostics         | pairwise_policy_robustness.csv               |
| Integrity check            | SHA256SUMS; claim audit JSON                 |

**Greedy matching pseudocode.** Sort retained predictions by descending score. For each prediction, form eligible unmatched GT boxes with IoU at least  $\iota$ ; if the set is nonempty, match its highest-IoU member and increment TP, otherwise increment FP. After all predictions, unmatched retained GT boxes increment FN. Input order is never used to resolve a metric tie: exact candidate ties are recorded as rank instability.

**Threshold-first Hungarian pseudocode.** Build the prediction–GT bipartite graph only from edges with IoU at least  $\iota$ . First maximize cardinality; among maximum-cardinality assignments, maximize summed IoU; use dummy nodes for unmatched objects. Count matched valid edges as TP, unmatched predictions as FP, and unmatched GT boxes as FN. The released test suite contains a counterexample for cardinality-first optimization, plus no-op policy, input-order, and prediction-order checks.




Fig. 1. Best observed source-level  $F_1$  versus confidence. Each point maximizes over registered support thresholds, primary policies, IoUs, and candidates available for that source. It is an optimistic diagnostic envelope, not a fixed-protocol performance curve. AI-TOD remains descriptive because it has one candidate and incomplete coverage.




Fig. 2. Best candidate  $F_1$  over registered support thresholds and primary policies at each confidence-IoU coordinate: absolute support (top) and normalized support (bottom). VisDrone and local UAVDT use common-coverage pools; AI-TOD has one candidate and is descriptive only.

TABLE V  
AP<sub>50</sub> PAIRED-BOOTSTRAP HEADLINE RECORDS. DIFFERENCES ARE WINNER MINUS RUNNER-UP.

| Source      | Rule       | Threshold | Policy      | Winner | Difference | 95% interval | Winner $p$ |
|-------------|------------|-----------|-------------|--------|------------|--------------|------------|
| VisDrone    | absolute   | 24 px     | include-all | SAHI   | .354       | [.339,.366]  | 1.000      |
| VisDrone    | absolute   | 24 px     | exclude/off | SAHI   | .373       | [.361,.385]  | 1.000      |
| VisDrone    | absolute   | 24 px     | exclude/on  | SAHI   | .373       | [.361,.385]  | 1.000      |
| VisDrone    | normalized | .015      | include-all | SAHI   | .354       | [.339,.366]  | 1.000      |
| VisDrone    | normalized | .015      | exclude/off | SAHI   | .375       | [.363,.386]  | 1.000      |
| VisDrone    | normalized | .015      | exclude/on  | SAHI   | .375       | [.362,.387]  | 1.000      |
| Local UAVDT | absolute   | 24 px     | include-all | tiling | .034       | [-.296,.108] | .585       |
| Local UAVDT | absolute   | 24 px     | exclude/off | tiling | .028       | [-.310,.102] | .577       |
| Local UAVDT | normalized | .015      | include-all | tiling | .034       | [-.303,.111] | .604       |
| Local UAVDT | normalized | .015      | exclude/off | tiling | .001       | [-.303,.071] | .471       |

| Claim group              | Derived artifact          | Records   |                       |
|--------------------------|---------------------------|-----------|-----------------------|
| Coverage/support         | coverage and scale audits | 4 sources | exact count           |
| AP/ $F_1$ policy results | headline tables and grids | 42/9,408  | numerical             |
| Matching control         | Hungarian agreement file  | 84        | winner/runner-up      |
| Uncertainty              | paired bootstrap tables   | 34        | seed and interval     |
| Detection cap            | maxDets controls          | 20        | truncation/generation |

TABLE VI  
COMPLETE MICRO- $F_1$  PAIRED-BOOTSTRAP HEADLINE RECORDS. DIFFERENCES ARE WINNER MINUS RUNNER-UP; EACH ROW USES 1,000 PAIRED REPLICATES.

| Source      | Rule       | Threshold | Policy      | IoU  | Winner | Difference | 95% interval   | Winner $p$ |
|-------------|------------|-----------|-------------|------|--------|------------|----------------|------------|
| VisDrone    | absolute   | 24 px     | include-all | 0.25 | SAHI   | 0.335      | [0.328,0.342]  | 1.000      |
| VisDrone    | absolute   | 24 px     | include-all | 0.50 | SAHI   | 0.253      | [0.241,0.264]  | 1.000      |
| VisDrone    | absolute   | 24 px     | exclude/off | 0.25 | SAHI   | 0.322      | [0.316,0.330]  | 1.000      |
| VisDrone    | absolute   | 24 px     | exclude/off | 0.50 | SAHI   | 0.274      | [0.265,0.284]  | 1.000      |
| VisDrone    | absolute   | 24 px     | exclude/on  | 0.25 | SAHI   | 0.323      | [0.316,0.330]  | 1.000      |
| VisDrone    | absolute   | 24 px     | exclude/on  | 0.50 | SAHI   | 0.275      | [0.265,0.284]  | 1.000      |
| VisDrone    | normalized | .015      | include-all | 0.25 | SAHI   | 0.335      | [0.328,0.342]  | 1.000      |
| VisDrone    | normalized | .015      | include-all | 0.50 | SAHI   | 0.253      | [0.241,0.264]  | 1.000      |
| VisDrone    | normalized | .015      | exclude/off | 0.25 | SAHI   | 0.331      | [0.324,0.338]  | 1.000      |
| VisDrone    | normalized | .015      | exclude/off | 0.50 | SAHI   | 0.277      | [0.267,0.287]  | 1.000      |
| VisDrone    | normalized | .015      | exclude/on  | 0.25 | SAHI   | 0.332      | [0.325,0.339]  | 1.000      |
| VisDrone    | normalized | .015      | exclude/on  | 0.50 | SAHI   | 0.277      | [0.268,0.286]  | 1.000      |
| Local UAVDT | absolute   | 24 px     | include-all | 0.25 | tiling | 0.191      | [-0.238,0.282] | 0.786      |
| Local UAVDT | absolute   | 24 px     | include-all | 0.50 | tiling | 0.162      | [-0.258,0.259] | 0.735      |
| Local UAVDT | absolute   | 24 px     | exclude/off | 0.25 | tiling | 0.185      | [-0.229,0.275] | 0.811      |
| Local UAVDT | absolute   | 24 px     | exclude/off | 0.50 | tiling | 0.156      | [-0.251,0.251] | 0.716      |
| Local UAVDT | absolute   | 24 px     | exclude/on  | 0.25 | tiling | 0.185      | [-0.243,0.275] | 0.749      |
| Local UAVDT | absolute   | 24 px     | exclude/on  | 0.50 | tiling | 0.156      | [-0.261,0.257] | 0.725      |
| Local UAVDT | normalized | .015      | include-all | 0.25 | tiling | 0.191      | [-0.238,0.281] | 0.808      |
| Local UAVDT | normalized | .015      | include-all | 0.50 | tiling | 0.162      | [-0.250,0.256] | 0.721      |
| Local UAVDT | normalized | .015      | exclude/off | 0.25 | tiling | 0.135      | [-0.231,0.218] | 0.775      |
| Local UAVDT | normalized | .015      | exclude/off | 0.50 | tiling | 0.108      | [-0.256,0.196] | 0.691      |
| Local UAVDT | normalized | .015      | exclude/on  | 0.25 | tiling | 0.135      | [-0.234,0.217] | 0.782      |
| Local UAVDT | normalized | .015      | exclude/on  | 0.50 | tiling | 0.108      | [-0.235,0.194] | 0.705      |

**Detection-cap audit.** The released truncation table records, for every multi-candidate AP setting, the number of images affected and detections removed at  $\text{maxDets} = 100, 300$ , and 2000. At 2000, all ten retained-prediction settings have zero truncation; the largest per-image retained prediction count is 1,066. The paired  $\text{AP}_{50}$  winner is unchanged in all ten nonduplicate outcomes, but the corresponding winner–runner-up gaps move by  $[-.041, .048]$  from 100 to 2000. This control is reported as evaluator sensitivity, not as a recommendation to replace the registered 100-detection COCO setting.

**Claim-audit summary.** The 55 verified numerical claims comprise artifact counts and support summaries (16), AP and  $F_1$  policy quantities (18), matching and tie-sensitive quantities (5), bootstrap endpoints/probabilities (10), and detection-cap controls (6). The audit checks the manuscript’s rounded values against named released derived artifacts; it does not certify the external licensed inputs. Each candidate manifest entry records a path, inclusion decision, and SHA-256 value, while  $\text{SHA256SUMS}$  validates the data-free supplementary repository files.

A minimal reproduction order is: verify hashes; inspect coverage and the frozen candidate manifest; evaluate AP and coverage-corrected absolute/normalized grids; run threshold-first Hungarian and removed-as-ignore controls; run paired bootstraps; regenerate the pairwise-policy diagnostic; and compile the manuscript. The data-free supplementary repository deliberately omits licensed raw imagery and prediction Parquet files while retaining their manifest paths and hashes. The

automated manuscript claim audit verifies 55 numerical claims against released derived records. This makes the scope of each figure, table, and headline statement inspectable without extending the declared candidate pool or rerunning model training.