

Supplementary Material for “Coverage-Qualified Evaluation-Contract Auditing for Aerial Vehicle Detection”

Xinling Wen, Jiahao Zhang, Wenqi Liu, Yangxu Ren, and Yu Chen

TABLE S1

DECLARED EVALUATION CONTRACTS. G IS ALL VALID GT, G_t THE SUPPORT-RETAINED TARGET, $D_t = G \setminus G_t$, AND H SOURCE IGNORE.

ID	Scored valid GT	Neutral regions	Step meaning
$A \ G$		none	all-valid reference
$F \ G_t$		none	removed objects are background
$R \ G_t$		D_t	neutralize removed objects
$S \ G_t$		$D_t \cup H$	add source ignores

TABLE S2

FIVE-CONFIGURATION VISDRONE PATH SUMMARY AT CONFIDENCE AND VALID-MATCH IOU .25. A = ABSOLUTE 24 PX; N = NORMALIZED .015.

Configuration	Rule	δ_{target}	δ_{removed}	δ_{source}	Range
B-640	A	-.0360	+.0667	+.0093	.0760
	N	+.0178	+.0093	+.0078	.0349
B-1280	A	-.0930	+.0911	+.0281	.1192
	N	-.0192	+.0252	+.0237	.0489
Y11m	A	-.0310	+.0727	+.0119	.0846
	N	+.0254	+.0088	+.0100	.0441
ASD	A	-.0759	+.0932	+.0166	.1099
	N	-.0025	+.0203	+.0138	.0341
Y11m-P2	A	-.0699	+.0922	+.0159	.1080
	N	-.0032	+.0216	+.0131	.0348

S1. DECLARED CONTRACTS AND COUNT LEDGER

The four contracts alter one evaluator factor at a time. All detections are score ordered; retained valid GT is matched before an otherwise unmatched detection can be neutralized. Removed and source-ignore regions use intersection over prediction area at .5. Table S1 distinguishes the scored target from neutral regions explicitly.

Table S2 summarizes the five support-conditioned count-transition ledgers behind Fig. 2 of the main paper. Every row comes from one frozen prediction artifact. The released 90-row ledger contains the corresponding integer TP/FP/FN changes. Each three-step sum equals $F_{1,S} - F_{1,A}$ to machine precision.

Table S3 extends the path beyond one baseline per source. VisDrone contributes five separately materialized artifacts. UAVDT contributes three resolution slices and a tiling configuration. AI-TOD contributes three resolution slices on the same covered subset; these rows test cache-resolution sensitivity without creating a full-benchmark estimand.

Removed-object neutralization is positive in all 24 configuration–rule replays. Source-ignore handling is positive in all ten VisDrone replays and exactly zero for UAVDT and AI-TOD, which have no active source-ignore regions. Target filtering remains configuration dependent, especially under normalized support. The exact 96-row endpoint record and

72-row transition ledger preserve artifact groups and input hashes.

S2. COVERAGE QUALIFICATION AND FULL SUPPORT GRID

Raw prediction names establish coverage before class filtering. VisDrone and UAVDT cover all declared images. AI-TOD covers 1,869/2,804 images. Its full-annotation support diagnosis uses all 2,804 images, while the ancillary metric replay is confined to the covered subset. The 1,869 covered images contain 16,900 boxes; the 935 noncovered images contain 43,004.

The image-level record shows that incomplete coverage is not exchangeable with the full validation set. The covered subset has many empty vehicle images, yet among nonempty images its per-image support composition also differs. At the box level, covered and noncovered median max sides are 12 and 16 pixels; the normalized-support Wasserstein distance is .00370. No weighting or imputation is used to extrapolate the conditional metric endpoints.

S3. PAIRED RANK RECORDS AND EVALUATOR CONTROLS

Table S6 gives the localization-by-metric control for the Y11m–ASD pair. Sequence clusters are the leading seven-digit VisDrone filename tokens (76 clusters). Every replicate includes all images in each sampled sequence, applies identical multiplicities to both configurations, and recomputes global counts or pooled-detection AP. The max- t critical value is 2.744 across the eight primary comparisons.

Fixed-confidence F_1 has positive pointwise intervals in all four cells and remains an operational secondary endpoint. The simultaneous family resolves three max- F_1 cells and two AP cells. Normalized max- F_1 at IoU .50 joins the two previously unresolved AP cells after multiplicity adjustment.

The alternative definitions preserve every max- F_1 and AP direction in the 2×2 design. Both neutralization steps remain positive in all 30 candidate–rule combinations.

The AP cap is 2,000 detections per image. The observed maximum is 1,881, so no pool image is truncated; caps of 100 and 300 alter dense-image inputs. The independent AP implementation and COCOeval agree exactly for AP50 (maximum absolute difference $< 10^{-12}$). AP25 keeps the S -contract neutralization threshold fixed at .5, whereas COCOeval couples crowd matching to the evaluated IoU; COCOeval equivalence is therefore not asserted for that deliberately controlled cell. A threshold-first Hungarian replay selects the

TABLE S3

CROSS-CONFIGURATION CONTRACT-PATH RANGES AT CONFIDENCE AND VALID-MATCH IOU .25. EACH INTERVAL SPANS THE CONFIGURATIONS WITHIN ONE SOURCE. A = 24 PX; N = .015. AI-TOD IS CONDITIONED ON 1,869/2,804 IMAGES.

Source	Rule	Configs.	δ_{target}	δ_{removed}	δ_{source}	Contract range
VisDrone	A	5	-.0930-+.0310	+.0667-+.0932	+.0093-+.0281	.0760-.1192
	N	5	-.0192-+.0254	+.0088-+.0252	+.0078-+.0237	.0341-.0489
UAVDT	A	4	-.0147-+.0001	+.0029-+.0067	0	.0030-.0147
	N	4	-.0041-+.0025	+.0010-+.0027	0	.0018-.0041
AI-TOD ^a	A	3	-.5105-+.4565	+.1450-+.1920	0	.4565-.5105
	N	3	-.1620-+.0943	+.0859-+.1028	0	.0943-.1620

^aCoverage-conditioned resolution control; all three slices share one checkpoint and covered image set.

TABLE S4

AI-TOD COVERED–NONCOVERED IMAGE QUALIFICATION. BRACKETS ARE MEDIAN [Q1,Q3]. SMD USES COVERED MINUS NONCOVERED MEANS; W_1 IS IN THE FEATURE’S NATIVE UNIT.

Feature	Covered	Noncovered	SMD	W_1
Boxes/image	0 [0,1]	21 [7,58.5]	-.673	36.951
Median side (px) ^a	17 [12,22]	16 [14,18]	+.448	5.597
Fraction < 24 px ^a	.895 [.600,1]	.961 [.857,1]	-.640	.159
Fraction < .015 ^a	.018 [0,.410]	.009 [0,.111]	+.610	.148
Nearest distance ^b	.031 [.016,.065]	.027 [.018,.055]	+.029	.009

^aNonempty images. ^bImages with at least two vehicle boxes; centers normalized by image size.

TABLE S5

COMPLETE POOLED RETAINED VALID VEHICLE GT (%) BEFORE PREDICTION MATCHING.

Rule	Threshold	VisDrone	UAVDT	AI-TOD
Absolute	16 px	85.7	99.6	47.6
	20 px	79.5	98.8	20.4
	24 px	73.3	95.1	9.7
	28 px	67.5	77.5	5.8
	32 px	62.1	56.4	4.2
	40 px	52.3	34.3	2.5
	48 px	42.7	30.0	1.4
	64 px	27.7	17.7	0.4
Normalized	.0050	100.0	100.0	99.9
	.0075	98.8	100.0	99.6
	.0100	96.5	99.8	98.2
	.0150	89.8	97.4	81.0
	.0200	82.4	80.0	47.6
	.0300	67.6	46.4	9.7

TABLE S6

Y11M–ASD DIFFERENCES UNDER TERMINAL CONTRACT S (10,000 PAIRED SEQUENCE-CLUSTER REPLICATES). A = 24 PX; N = .015.

Rule	IoU/metric	Diff.	Pointwise 95%	Simultaneous 95%	p_{adj}
A	.25/max- F_1	+.0163	[.0119,.0213]	[.0097,.0230]	< .001
	.25/AP	+.0081	[.0051,.0120]	[.0033,.0129]	< .001
	.50/max- F_1	+.0141	[.0099,.0191]	[.0076,.0207]	< .001
	.50/AP	+.0013	[-.0033,.0052]	[-.0046,.0072]	.984
N	.25/max- F_1	+.0099	[.0060,.0140]	[.0043,.0155]	< .001
	.25/AP	-.0005	[-.0030,.0047]	[-.0061,.0051]	1.000
	.50/max- F_1	+.0054	[.0013,.0103]	[-.0010,.0117]	.125
	.50/AP	-.0110	[-.0181,-.0057]	[-.0199,-.0022]	.008

same top configuration as greedy matching in 24/24 settings; the maximum absolute F_1 difference is .00546.

The full 14-support-by-eight-confidence rank surface remains exploratory because its small-gap cells were identified after inspecting the prespecified grid. No pointwise interval from a selected boundary cell is used as confirmatory evidence. Machine-readable full surfaces, matching controls, maxDets diagnostics, and DOTA-v2 structural support distances remain in the fixed repository release.

S4. ANCILLARY QUALIFICATION CONTROLS

Figure 1 separates two boundary checks that precede metric interpretation. The AI-TOD density gap is visible at the image level, while the box-level curves show that coverage selection also changes the support composition of the observed target. The cap panel uses mapped vehicle predictions before AP truncation: a low cap would modify the prediction set for every configuration by a different amount. The declared cap of 2,000 exceeds the observed per-image maximum of 1,881 and therefore preserves every mapped prediction.

Figure 2 provides the structural context omitted from the main figures. Normalization changes the coordinate system and greatly compresses the numerical scale, but it does not collapse the four empirical distributions. In particular, the largest normalized pairwise distance is .031 between VisDrone and AI-TOD, whereas VisDrone and DOTA-v2 differ by .006. These are distributional distances, not annotation-quality scores, and no detector result is inferred from them.

S5. REPRODUCIBILITY RECORD

The data-free release provides configurations, derived records, scripts, and SHA-256 checksums. Licensed imagery, trained weights, and raw prediction caches remain external and are referenced by manifest path and hash. Reproduction order is coverage → support → contract endpoints/ledger → paired metric qualification → figures.

TABLE S7

RETENTION-MATCHED SUPPORT-DEFINITION CONTROLS. RANGE IS THE SPAN OF FIVE CANDIDATE-LEVEL $A \rightarrow F \rightarrow R \rightarrow S$ RANGES. REMAINING COLUMNS ARE Y11m-ASD TERMINAL- S POINT DIFFERENCES.

Coordinate	Support definition (threshold)	Retained	Path range	max- $F_{1,25}$	AP25	max- $F_{1,50}$	AP50
Absolute	max side (24 px)	73.33%	.076-.119	+.0163	+.0081	+.0141	+.0013
	minimum side (17 px)	73.95%	.067-.115	+.0174	+.0072	+.0152	+.0019
	geometric mean (20.149 px)	73.30%	.073-.119	+.0178	+.0083	+.0161	+.0025
Normalized	max side (.015)	89.77%	.034-.049	+.0099	-.0005	+.0054	-.0110
	minimum side (.009559)	90.32%	.032-.047	+.0088	-.0025	+.0045	-.0157
	geometric mean (.012187)	89.77%	.034-.049	+.0103	-.0009	+.0058	-.0113




Fig. 1. Ancillary qualification records. (a)–(c) AI-TOD covered and noncovered image groups differ in density and support; vertical lines mark 24 px and .015. ECDFs in (b) and (c) use individual vehicle boxes. (d) Fractions of the 548 common VisDrone images whose mapped prediction counts exceed maxDets 100 or 300. At maxDets 2000, all five fractions are zero.




Fig. 2. Pairwise empirical Wasserstein distances for full-annotation vehicle support. (a) Absolute max side is scale covariant. (b) Normalized max side is resize invariant but retains nonzero source distances. DOTA-v2 is structural-only evidence and does not enter prediction-conditioned metric or rank claims.




Fig. 3. Cross-source summary of the declared headline path. (a) Four-contract micro- F_1 ranges. (b) Adjacent path changes, preserving the declared order rather than treating them as order-free causal components. Normalized .015 reduces the range in all three sources but leaves source-dependent residual sensitivity.

TABLE S8
MINIMAL REPLAY CHAIN IN THE RELEASED REPOSITORY.

Audit object	Released record or script
Coverage gate	outputs/coverage_qualification/ and evaluate_aitod_coverage_qualification.py
Support denomina- tor	outputs/scale/box_scale_records.parquet and raw-coverage manifests
Contract path	outputs/contract_path/, contract_path.py, and evaluate_contract_path_headlines.py
Metric-qualified rank	outputs/metric_qualified_rank/ and evaluate_metric_qualified_visdrone_rank.py
Metric-factorial control	outputs/metric_factorial_control/ and evaluate_metric_factorial_control.py
Figures and tests	Four build_*.py scripts; contract, support, matching, and AP regression tests
Repository	https://github.com/Marchemantics/aerial-evaluation-audit